

# The Java 8 Stream API Tutorial

 [baeldung.com/java-8-streams](https://baeldung.com/java-8-streams)

By baeldung

June 15, 2016



## 1. Overview

In this comprehensive tutorial, we'll go through the practical uses of Java 8 Streams from creation to parallel execution.

To understand this material, readers need to have a basic knowledge of Java 8 (lambda expressions, *Optional*, method references) and of the Stream API. In order to be more familiar with these topics, please take a look at our previous articles: [New Features in Java 8](#) and [Introduction to Java 8 Streams](#).

## 2. Stream Creation

There are many ways to create a stream instance of different sources. Once created, the instance **will not modify its source**, therefore allowing the creation of multiple instances from a single source.

### 2.1. Empty Stream

We should use the ***empty()*** method in case of the creation of an empty stream:

```
Stream<String> streamEmpty = Stream.empty();
```

We often use the ***empty()*** method upon creation to avoid returning for streams with no element:

```
public Stream<String> streamOf(List<String> list) {  
    return list == null || list.isEmpty() ? Stream.empty() : list.stream();  
}
```

## 2.2. Stream of Collection

---

We can also create a stream of any type of *Collection* (*Collection*, *List*, *Set*):

```
Collection<String> collection = Arrays.asList("a", "b", "c");
Stream<String> streamOfCollection = collection.stream();
```

## 2.3. Stream of Array

---

An array can also be the source of a stream:

```
Stream<String> streamOfArray = Stream.of("a", "b", "c");
```

We can also create a stream out of an existing array or of part of an array:

```
String[] arr = new String[]{"a", "b", "c"};
Stream<String> streamOfArrayFull = Arrays.stream(arr);
Stream<String> streamOfArrayPart = Arrays.stream(arr, 1, 3);
```

## 2.4. Stream.builder()

---

When **builder** is used, the desired type should be additionally specified in the right part of the statement, otherwise the *build()* method will create an instance of the *Stream<Object>*:

```
Stream<String> streamBuilder =
    Stream.<String>builder().add("a").add("b").add("c").build();
```

## 2.5. Stream.generate()

---

The **generate()** method accepts a *Supplier<T>* for element generation. As the resulting stream is infinite, the developer should specify the desired size, or the *generate()* method will work until it reaches the memory limit:

```
Stream<String> streamGenerated =
    Stream.generate(() -> "element").limit(10);
```

The code above creates a sequence of ten strings with the value “*element.*”

## 2.6. Stream.iterate()

---

Another way of creating an infinite stream is by using the **iterate()** method:

```
Stream<Integer> streamIterated = Stream.iterate(40, n -> n + 2).limit(20);
```

The first element of the resulting stream is the first parameter of the *iterate()* method. When creating every following element, the specified function is applied to the previous element. In the example above the second element will be 42.

## 2.7. Stream of Primitives

---

Java 8 offers the possibility to create streams out of three primitive types: *int*, *long* and *double*. As *Stream<T>* is a generic interface, and there is no way to use primitives as a type parameter with generics, three new special interfaces were created: ***IntStream***, ***LongStream***, ***DoubleStream***.

Using the new interfaces alleviates unnecessary auto-boxing, which allows for increased productivity:

```
IntStream intStream = IntStream.range(1, 3);
LongStream longStream = LongStream.rangeClosed(1, 3);
```

The ***range(int startInclusive, int endExclusive)*** method creates an ordered stream from the first parameter to the second parameter. It increments the value of subsequent elements with the step equal to 1. The result doesn't include the last parameter, it is just an upper bound of the sequence.

The ***rangeClosed(int startInclusive, int endInclusive)*** method does the same thing with only one difference, the second element is included. We can use these two methods to generate any of the three types of streams of primitives.

Since Java 8, the *Random* class provides a wide range of methods for generating streams of primitives. For example, the following code creates a *DoubleStream*, which has three elements:

```
Random random = new Random();
DoubleStream doubleStream = random.doubles(3);
```

## 2.8. Stream of *String*

---

We can also use *String* as a source for creating a stream with the help of the *chars()* method of the *String* class. Since there is no interface for *CharStream* in JDK, we use the *IntStream* to represent a stream of chars instead.

```
IntStream streamOfChars = "abc".chars();
```

The following example breaks a *String* into sub-strings according to specified *RegEx*:

```
Stream<String> streamOfString =
    Pattern.compile(", ").splitAsStream("a, b, c");
```

## 2.9. Stream of File

---

Furthermore, Java NIO class *Files* allows us to generate a *Stream<String>* of a text file through the *lines()* method. Every line of the text becomes an element of the stream:

```
Path path = Paths.get("C:\\file.txt");
Stream<String> streamOfStrings = Files.lines(path);
Stream<String> streamWithCharset =
    Files.lines(path, Charset.forName("UTF-8"));
```

The *Charset* can be specified as an argument of the *lines()* method.

### 3. Referencing a Stream

---

We can instantiate a stream, and have an accessible reference to it, as long as only intermediate operations are called. Executing a terminal operation makes a stream inaccessible.

To demonstrate this, we will forget for a while that the best practice is to chain the sequence of operation. Besides its unnecessary verbosity, technically the following code is valid:

```
Stream<String> stream =  
    Stream.of("a", "b", "c").filter(element -> element.contains("b"));  
Optional<String> anyElement = stream.findAny();
```

However, an attempt to reuse the same reference after calling the terminal operation will trigger the *IllegalStateException*:

```
Optional<String> firstElement = stream.findFirst();
```

As the *IllegalStateException* is a *RuntimeException*, a compiler will not signalize about a problem. So it is very important to remember that **Java 8 streams can't be reused**.

This kind of behavior is logical. We designed streams to apply a finite sequence of operations to the source of elements in a functional style, not to store elements.

So to make the previous code work properly, some changes should be made:

```
List<String> elements =  
    Stream.of("a", "b", "c").filter(element -> element.contains("b"))  
        .collect(Collectors.toList());  
Optional<String> anyElement = elements.stream().findAny();  
Optional<String> firstElement = elements.stream().findFirst();
```

### 4. Stream Pipeline

---

To perform a sequence of operations over the elements of the data source and aggregate their results, we need three parts: the **source**, **intermediate operation(s)** and a **terminal operation**.

Intermediate operations return a new modified stream. For example, to create a new stream of the existing one without few elements, the *skip()* method should be used:

```
Stream<String> onceModifiedStream =  
    Stream.of("abcd", "bbcd", "cbcd").skip(1);
```

If we need more than one modification, we can chain intermediate operations. Let's assume that we also need to substitute every element of the current *Stream<String>* with a sub-string of the first few chars. We can do this by chaining the *skip()* and *map()* methods:

```
Stream<String> twiceModifiedStream =  
    stream.skip(1).map(element -> element.substring(0, 3));
```

As we can see, the *map()* method takes a lambda expression as a parameter. If we want to learn more about lambdas, we can take a look at our tutorial [Lambda Expressions and Functional Interfaces: Tips and Best Practices](#).

A stream by itself is worthless; the user is interested in the result of the terminal operation, which can be a value of some type or an action applied to every element of the stream. **We can only use one terminal operation per stream.**

The correct and most convenient way to use streams is by a **stream pipeline, which is a chain of the stream source, intermediate operations, and a terminal operation:**

```
List<String> list = Arrays.asList("abc1", "abc2", "abc3");  
long size = list.stream().skip(1)  
    .map(element -> element.substring(0, 3)).sorted().count();
```

## 5. Lazy Invocation

---

**Intermediate operations are lazy.** This means that **they will be invoked only if it is necessary for the terminal operation execution.**

For example, let's call the method *wasCalled()*, which increments an inner counter every time it's called:

```
private long counter;  
  
private void wasCalled() {  
    counter++;  
}
```

Now let's call the method *wasCalled()* from operation *filter()*:

```
List<String> list = Arrays.asList("abc1", "abc2", "abc3");  
counter = 0;  
Stream<String> stream = list.stream().filter(element -> {  
    wasCalled();  
    return element.contains("2");  
});
```

As we have a source of three elements, we can assume that the *filter()* method will be called three times, and the value of the *counter* variable will be 3. However, running this code doesn't change *counter* at all, it is still zero, so the *filter()* method wasn't even called once. The reason why is missing of the terminal operation.

Let's rewrite this code a little bit by adding a *map()* operation and a terminal operation, *findFirst()*. We will also add the ability to track the order of method calls with the help of logging:

```
Optional<String> stream = list.stream().filter(element -> {
    log.info("filter() was called");
    return element.contains("2");
}).map(element -> {
    log.info("map() was called");
    return element.toUpperCase();
}).findFirst();
```

The resulting log shows that we called the *filter()* method twice and the *map()* method once. This is because the pipeline executes vertically. In our example, the first element of the stream didn't satisfy the filter's predicate. Then we invoked the *filter()* method for the second element, which passed the filter. Without calling the *filter()* for the third element, we went down through the pipeline to the *map()* method.

The *findFirst()* operation satisfies by just one element. So in this particular example, the lazy invocation allowed us to avoid two method calls, one for the *filter()* and one for the *map()*.

## 6. Order of Execution

---

From the performance point of view, **the right order is one of the most important aspects of chaining operations in the stream pipeline:**

```
long size = list.stream().map(element -> {
    wasCalled();
    return element.substring(0, 3);
}).skip(2).count();
```

Execution of this code will increase the value of the counter by three. This means that we called the *map()* method of the stream three times, but the value of the *size* is one. So the resulting stream has just one element, and we executed the expensive *map()* operations for no reason two out of the three times.

If we change the order of the *skip()* and the *map()* methods, the *counter* will increase by only one. So we will call the *map()* method only once:

```
long size = list.stream().skip(2).map(element -> {
    wasCalled();
    return element.substring(0, 3);
}).count();
```

This brings us to the following rule: **intermediate operations which reduce the size of the stream should be placed before operations which are applying to each element.** So we need to keep methods such as *skip()*, *filter()*, and *distinct()* at the top of our stream pipeline.

## 7. Stream Reduction

---

The API has many terminal operations which aggregate a stream to a type or to a primitive: *count()*, *max()*, *min()*, and *sum()*. However, these operations work according to the predefined implementation. So what **if a developer needs to customize a Stream's reduction mechanism?** There are two methods which allow us to do this, the *reduce()* and the *collect()* methods.

## 7.1. The *reduce()* Method

---

There are three variations of this method, which differ by their signatures and returning types. They can have the following parameters:

**identity** – the initial value for an accumulator, or a default value if a stream is empty and there is nothing to accumulate

**accumulator** – a function which specifies the logic of the aggregation of elements. As the accumulator creates a new value for every step of reducing, the quantity of new values equals the stream's size and only the last value is useful. This is not very good for the performance.

**combiner** – a function which aggregates the results of the accumulator. We only call combiner in a parallel mode to reduce the results of accumulators from different threads.

Now let's look at these three methods in action:

```
OptionalInt reduced =  
    IntStream.range(1, 4).reduce((a, b) -> a + b);
```

*reduced* = 6 (1 + 2 + 3)

```
int reducedTwoParams =  
    IntStream.range(1, 4).reduce(10, (a, b) -> a + b);
```

*reducedTwoParams* = 16 (10 + 1 + 2 + 3)

```
int reducedParams = Stream.of(1, 2, 3)  
    .reduce(10, (a, b) -> a + b, (a, b) -> {  
        log.info("combiner was called");  
        return a + b;  
    });
```

The result will be the same as in the previous example (16), and there will be no log in, which means that combiner wasn't called. To make a combiner work, a stream should be parallel:

```
int reducedParallel = Arrays.asList(1, 2, 3).parallelStream()  
    .reduce(10, (a, b) -> a + b, (a, b) -> {  
        log.info("combiner was called");  
        return a + b;  
    });
```

The result here is different (36), and the combiner was called twice. Here the reduction works by the following algorithm: the accumulator ran three times by adding every element of the stream to *identity*. These actions are being done in parallel. As a result, they have ( $10 + 1 = 11$ ;  $10 + 2 = 12$ ;  $10 + 3 = 13$ ). Now combiner can merge these three results. It needs two iterations for that ( $12 + 13 = 25$ ;  $25 + 11 = 36$ ).

## 7.2. The `collect()` Method

---

The reduction of a stream can also be executed by another terminal operation, the `collect()` method. It accepts an argument of the type *Collector*, which specifies the mechanism of reduction. There are already created, predefined collectors for most common operations. They can be accessed with the help of the *Collectors* type.

In this section, we will use the following *List* as a source for all streams:

```
List<Product> productList = Arrays.asList(new Product(23, "potatoes"),
    new Product(14, "orange"), new Product(13, "lemon"),
    new Product(23, "bread"), new Product(13, "sugar"));
```

### Converting a stream to the *Collection* (*Collection*, *List* or *Set*):

```
List<String> collectorCollection =
    productList.stream().map(Product::getName).collect(Collectors.toList());
```

### Reducing to *String*:

```
String listToString = productList.stream().map(Product::getName)
    .collect(Collectors.joining(", ", "[", "]"));
```

The *joiner()* method can have from one to three parameters (delimiter, prefix, suffix). The most convenient thing about using *joiner()* is that the developer doesn't need to check if the stream reaches its end to apply the suffix and not to apply a delimiter. *Collector* will take care of that.

### Processing the average value of all numeric elements of the stream:

```
double averagePrice = productList.stream()
    .collect(Collectors.averagingInt(Product::getPrice));
```

### Processing the sum of all numeric elements of the stream:

```
int summingPrice = productList.stream()
    .collect(Collectors.summingInt(Product::getPrice));
```

The methods *averagingXX()*, *summingXX()* and *summarizingXX()* can work with primitives (*int*, *long*, *double*) and with their wrapper classes (*Integer*, *Long*, *Double*). One more powerful feature of these methods is providing the mapping. As a result, the developer doesn't need to use an additional *map()* operation before the *collect()* method.

### Collecting statistical information about stream's elements:



```
IntSummaryStatistics statistics = productList.stream()
    .collect(Collectors.summarizingInt(Product::getPrice));
```

By using the resulting instance of type *IntSummaryStatistics*, the developer can create a statistical report by applying the *toString()* method. The result will be a *String* common to this one “*IntSummaryStatistics{count=5, sum=86, min=13, average=17,200000, max=23}*.”

It is also easy to extract from this object separate values for *count*, *sum*, *min*, and *average* by applying the methods *getCount()*, *getSum()*, *getMin()*, *getAverage()*, and *getMax()*. All of these values can be extracted from a single pipeline.

### Grouping of stream’s elements according to the specified function:

```
Map<Integer, List<Product>> collectorMapOfLists = productList.stream()
    .collect(Collectors.groupingBy(Product::getPrice));
```

In the example above, the stream was reduced to the *Map*, which groups all products by their price.

### Dividing stream’s elements into groups according to some predicate:

```
Map<Boolean, List<Product>> mapPartitioned = productList.stream()
    .collect(Collectors.partitioningBy(element -> element.getPrice() > 15));
```

### Pushing the collector to perform additional transformation:

```
Set<Product> unmodifiableSet = productList.stream()
    .collect(Collectors.collectingAndThen(Collectors.toSet(),
        Collections::unmodifiableSet));
```

In this particular case, the collector has converted a stream to a *Set*, and then created the unchangeable *Set* out of it.

### Custom collector:

If for some reason a custom collector should be created, the easiest and least verbose way of doing so is to use the method *of()* of the type *Collector*.

```
Collector<Product, ?, LinkedList<Product>> toLinkedList =
    Collector.of(LinkedList::new, LinkedList::add,
        (first, second) -> {
            first.addAll(second);
            return first;
        });

LinkedList<Product> linkedListOfPersons =
    productList.stream().collect(toLinkedList);
```

In this example, an instance of the *Collector* got reduced to the *LinkedList<Person>*.

## 8. Parallel Streams

---

Before Java 8, parallelization was complex. The emergence of the [ExecutorService](#) and the [ForkJoin](#) simplified a developer's life a little bit, but it was still worth remembering how to create a specific executor, how to run it, and so on. Java 8 introduced a way of accomplishing parallelism in a functional style.

The API allows us to create parallel streams, which perform operations in a parallel mode. When the source of a stream is a *Collection* or an *array*, it can be achieved with the help of the ***parallelStream()*** method:

```
Stream<Product> streamOfCollection = productList.parallelStream();
boolean isParallel = streamOfCollection.isParallel();
boolean bigPrice = streamOfCollection
    .map(product -> product.getPrice() * 12)
    .anyMatch(price -> price > 200);
```

If the source of a stream is something other than a *Collection* or an *array*, the ***parallel()*** method should be used:

```
IntStream intStreamParallel = IntStream.range(1, 150).parallel();
boolean isParallel = intStreamParallel.isParallel();
```

Under the hood, Stream API automatically uses the *ForkJoin* framework to execute operations in parallel. By default, the common thread pool will be used and there is no way (at least for now) to assign some custom thread pool to it. [This can be overcome by using a custom set of parallel collectors.](#)

When using streams in parallel mode, avoid blocking operations. It is also best to use parallel mode when tasks need a similar amount of time to execute. If one task lasts much longer than the other, it can slow down the complete app's workflow.

The stream in parallel mode can be converted back to the sequential mode by using the ***sequential()*** method:

```
IntStream intStreamSequential = intStreamParallel.sequential();
boolean isParallel = intStreamSequential.isParallel();
```

## 9. Conclusion

---

The Stream API is a powerful, but simple to understand set of tools for processing the sequence of elements. When used properly, it allows us to reduce a huge amount of boilerplate code, create more readable programs, and improve an app's productivity.

In most of the code samples shown in this article, we left the streams unconsumed (we didn't apply the *close()* method or a terminal operation). In a real app, **don't leave an instantiated stream unconsumed, as that will lead to memory leaks.**

The complete code samples that accompany this article are available [over on GitHub.](#)